

A Review of Adversarial Neural Cryptography

Alisamar Husain

Dept. of Electrical Engineering
Jamia Millia Islamia

Abstract

Artificial neural networks are well known for their ability to selectively explore the solution space of a given problem. One of the recent applications of this feature is in the field of neural cryptography, which provides an opportunity to use ANNs to encrypt data such that it cannot be decrypted by an attacker.

In this paper we examine the efficacy, feasibility and general practicality of the use of *adversarial neural cryptography*, as coined by Abadi et al. in [?], and neural cryptography in general. We test systems recommended in the literature and examine their use in data transmission systems for the purpose of encrypting data from the perspective of securing a communication channel.

1 Introduction

The field of cryptography is broadly concerned with algorithms and protocols that ensure the secrecy and integrity of information. Cryptographic mechanisms are typically described as programs or Turing machines. By this definition, an appropriate neural network can possibly be considered a cryptographic function.

As neural networks are applied to increasingly complex tasks, they are often trained to meet end-to-end objectives that go beyond simple functional specifications. These objectives include, for example, generating realistic images [?] and solving multiagent problems.

Cryptography is broadly concerned with algorithms and protocols that ensure the secrecy and integrity of information. Cryptographic mechanisms are typically described as programs or Turing machines. Attackers are also described in those terms, with bounds on their complexity and on their chances of success. A mechanism is deemed secure if it achieves its goal against all attackers. For instance, an encryption algorithm is said to be secure if no attacker can extract information about plaintexts from ciphertexts. Modern cryptography provides rigorous versions of such definitions.

Neural networks are generally not meant to be great at cryptography. Famously, the simplest neural networks cannot even compute XOR, which is basic to many cryptographic algorithms. Nevertheless, as was demonstrated in the seminal paper on this subject [?] by Abadi et. al in 2016, neural networks can learn to protect the confidentiality of their data from other neural networks: they discover forms of encryption and decryption, without being taught specific algorithms for these purposes.

Knowing how to encrypt is seldom enough for security and privacy. Interestingly, neural networks can also learn what to encrypt in order to achieve a desired secrecy property while maximizing utility. Thus, when we wish to prevent an adversary from seeing a fragment of a plaintext, or from estimating a function of the plaintext, encryption can be selective, hiding the plaintext only partly.

1.1 Terminology

Certain terms are frequently used while talking about cryptographic mechanisms and it is beneficial to have an understanding of what these refer to. Some of these will be used in this paper to commonly identify certain parts of the system and some are abbreviations made for convenience.

A **party** is a machine, or actor in general, which is using a communication channel to communicate with another machine. There are two major types of parties which we are concerned with, **participants and attackers**.

A **participant** is a party which actively takes part in the communication and sends messages on the channel. The goal of encryption is to ensure that the communication between any two parties can only be intercepted and understood by them.

An **attacker** is a party which attempts to intercept and understand the communication between two participants.

Attackers Attackers are also described in those terms, with bounds on their complexity (e.g., limited to polynomial time) and on their chances of success (e.g., limited to a negligible probability). A mechanism is deemed secure if it achieves its goal against all attackers. For instance, an encryption algorithm is said to be secure if no attacker can extract information about plaintexts from ciphertexts. Modern cryptography provides rigorous versions of such definitions, like those given by Goldwasser & Micali. [?]

1.2 Symmetric Encryption

Symmetric-key cryptography refers to encryption methods in which both the sender and receiver share the same key (or, less commonly, in which their keys are different, but related in an easily computable way).

Symmetric key ciphers are implemented as either block ciphers or stream ciphers. A block cipher enciphers input in blocks of plaintext as opposed to individual characters, the input form used by a stream cipher.

1.3 Adversarial Neural Cryptography

This section discusses how to protect the confidentiality of plaintexts using shared keys. It describes the organization of the system that we consider, and the objectives of the participants in this system. It also explains the training of these participants, defines their architecture, and presents experiments.

A classic scenario in security involves three parties: Alice, Bob, and Eve. Typically, Alice and Bob wish to communicate securely, and Eve wishes to eavesdrop on their communications. Thus, the desired security property is secrecy (not integrity), and the adversary is a “passive attacker” that can intercept communications but that is otherwise quite limited: it cannot initiate sessions, inject messages, or modify messages in transit.

1.3.1 Mathematical Background

We write $A(\Phi_A, P, K)$ for Alice’s output on input P, K , write $B(\Phi_B, C, K)$ for Bob’s output on input C, K , and write $E(\Phi_E, C)$ for Eve’s output on input C . We introduce a distance function d on plaintexts.

Although the exact choice of this function is probably not crucial, for concreteness we take the L1 distance

$$d(P, P') = \sum_{i=1}^N |P_i - P'_i|$$

where N is the length of plaintexts. We define a per-example loss function for Eve. Intuitively this represents how much Eve is wrong when the plaintext is P and the key is K . We also define a loss function for Eve over the distribution on plaintexts and keys by taking an expected value:

$$L_E(\Phi_A, \Phi_E) = \mathbb{E}_{P, K} [d(P, E(\Phi_E, A(\Phi_A, P, K)))]$$

We obtain the “optimal Eve” by minimizing this loss:

$$O_E(\theta_A) = \operatorname{argmin}_{\theta_E} (L_E(\theta_A, \theta_E))$$

Similarly, we define a per-example reconstruction error for Bob, and extend it to the distribution on plaintexts and keys. We define a loss function for Alice and Bob by combining L_B and the optimal value of L_E :

$$L_{AB}(\Phi_A, \Phi_E) = L_B(\Phi_A, \Phi_B) - L_E(\Phi_A, O_E(\theta_A))$$

This combination reflects that Alice and Bob want to minimize Bob’s reconstruction error and to maximize the reconstruction error of the “optimal Eve”. The use of a simple subtraction is somewhat arbitrary; below we describe useful variants. We obtain the “optimal Alice and Bob” by minimizing

$$(O_A, O_B) = \operatorname{argmin}_{(\theta_A, \theta_B)} (L_{AB}(\theta_A, \theta_B))$$

We write “optimal” in quotes because there need be no single global minimum. In general, there are many equi-optimal solutions for Alice and Bob. As a simple example, assuming that the key is of the same size as the plaintext and the ciphertext, Alice and Bob may XOR the plaintext and the ciphertext, respectively, with any permutation of the key, and all permutations are equally good as long as Alice and Bob use the same one; moreover, with the way we architect our networks (see Section 2.4), all permutations are equally likely to arise.

2 Related Work

3 Methodology

We use a simple setup in order to build and test the networks. The models are first implemented using the Python programming language and after obtaining a suitably trained and validated model, we can move on to testing.

3.1 Network Architecture

The Architecture of Alice, Bob, and Eve Because we wish to explore whether a general neural network can learn to communicate securely, rather than to engineer a particular method, we aimed to create a neural network architecture that was sufficient to learn mixing functions such as XOR, but that did not strongly encode the form of any particular algorithm. To this end, we chose the following “mix & transform” architecture. It has a first fully-connected (FC) layer, where the number of outputs is equal to the number of inputs. The plaintext and key bits are fed into this FC layer. Because each output bit can be a linear combination of all of the input bits, this layer enables—but does not mandate—mixing between the key and the plaintext bits. In particular, this layer can permute the bits.

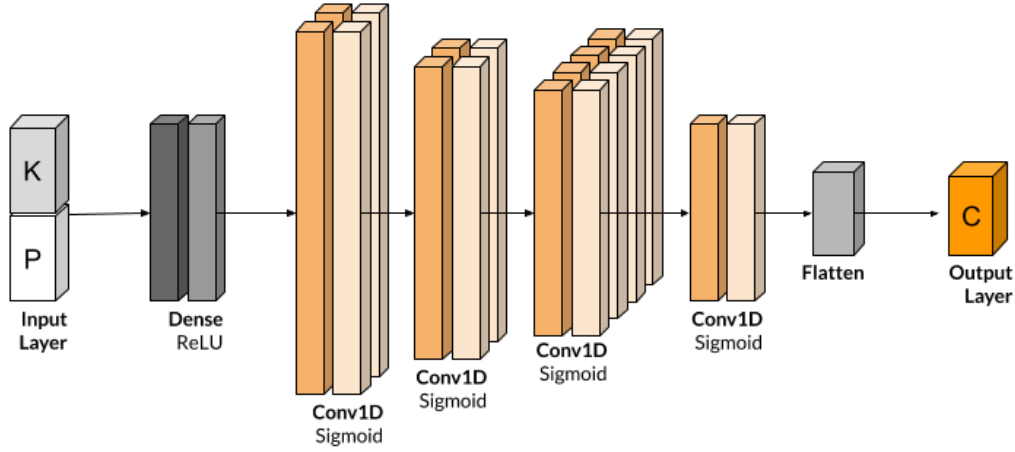


Figure 1: Modified network Model from Abadi et al.

The FC layer is followed by a sequence of convolutional layers, the last of which produces an output of a size suitable for a plaintext or ciphertext. These convolutional layers learn to apply some function to groups of the bits mixed by the previous layer, without an a priori specification of what that function should be. Notably, the opposite order (convolutional followed by FC) is much more common in image-processing applications. Neural networks developed for those applications frequently use convolutions to take advantage of spatial locality.

For neural cryptography, we specifically wanted locality—i.e., which bits to combine—to be a learned property, instead of a pre-specified one. While it would certainly work to manually pair each input plaintext bit with a corresponding key bit, we felt that doing so would be uninteresting. We refrain from imposing further constraints that would simplify the problem. For example, we do not tie the parameters Φ_A and Φ_B , as we would if we had in mind that Alice and Bob should both learn the same function, such as XOR.

3.2 Training Parameters

The parameters which aren’t learnt by the network but directly affect the performance are called **hyperparameters**. These are important since the right choice of hyperparameters can get us to the optimal model very quickly.

For our purposes the important parameters are the **key length**, **learning rate** and the **loss coefficients** in calculation of the loss for Alice and Bob. Various combinations are tried until one good set of values is obtained by experimentation.

Key Size	Learning Rate	Weight Decay
4 bits	0.0002	0.000001
8 bits	0.0004	0.000002
16 bits	0.0008	
32 bits		

3.3 Tools Used

The models are implemented using PyTorch, a popular framework for building neural networks in Python. The models are built per the specification given in the literature and trained with the help of several common Python libraries. We use Tensorboard to monitor the training process and log the training metadata.

4 Results

After training the networks suitably, they are put through the validation procedure to test their efficacy with unseen new data. The models that meet a reasonable level of validation accuracy are considered to have converged and thus a success.

4.1 ANC

The following figures show the loss trends for the ANC model with and without the adversary. We can see the loss decreasing fast without the adversary but when the eavesdropper is introduced, the loss still decreases. This is a sign that the cryptosystem is working [?].




Figure 2: Training Loss Trend with ANC

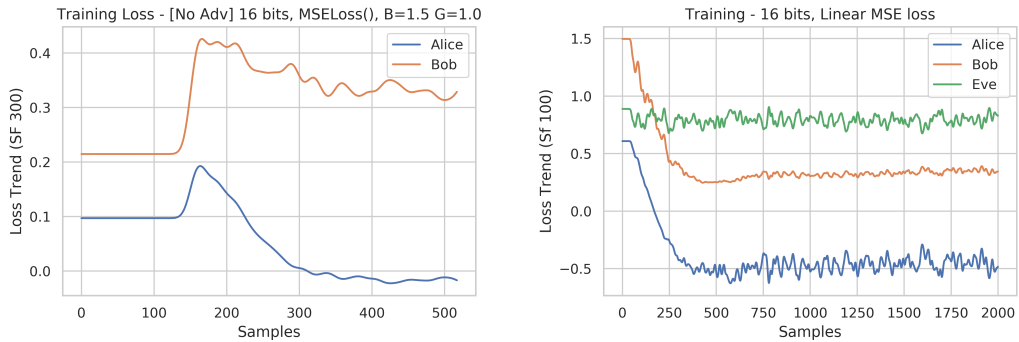


Figure 3: Loss Trend with ANC model by Abadi et al.

Key Size	Learning Rate	Epochs	Convergence
8 bits	0.0008	8	4/10
16 bits	0.0008	10	4/10
16 bits	0.0010	16	6/10

The table above shows the number of times out of 10, that the networks converged with a reasonable validation accuracy with different hyperparameters in the training.

4.2 Cryptonet

The following figures show the loss trends for the Cryptonet model with and without the adversary. Immediately we can notice that the rate of decrease of the loss is slower, which is a clear sign that the adversary Eve is having an effect on the training of Alice and Bob.




Figure 4: Training Loss Trend with Cryptonet

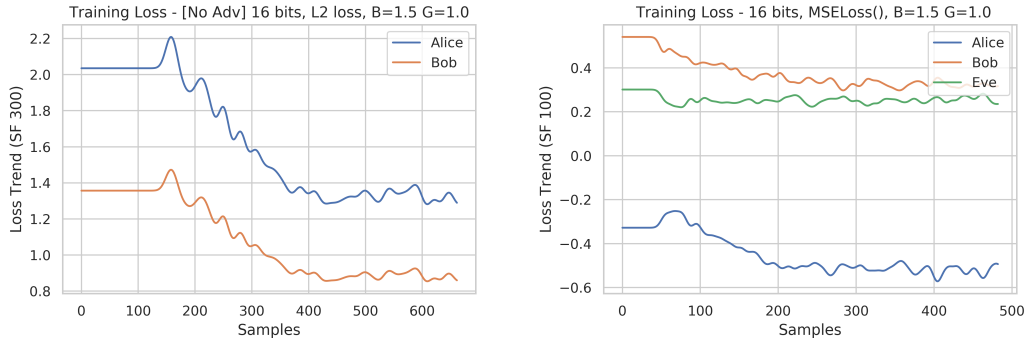


Figure 5: Loss Trend with Cryptonet model by Coutinho et al.

Key Size	Learning Rate	Epochs	Convergence
4 bits	0.0008	8	2/10
8 bits	0.0008	10	4/10
8 bits	0.0008	16	5/10

5 Conclusions